

LensaRak — Demo Video Script

4-Person Team | AI Innovators Challenge 2026

Target length: ≤ 5:00, ONE continuous take, unedited. Must show: the bulk-return scan, and at least one failure case with recovery.

Roles

Role	Who does what
R1 — Presenter / Narrator	Talks to camera throughout, introduces the team and the problem, narrates what's happening on-screen
R2 — Storeroom Operator	Physically places/returns items on the table, performs the deliberate low-confidence moment
R3 — System Operator	Sits at the laptop, runs LensaRak.bat, presses SPACE/keys, runs the review tool
R4 — Camera Operator	Holds the recording phone, frames the shot so BOTH the laptop screen and the item table are visible throughout

Setup note: R4's phone is the *recording* camera (for the video). A separate phone or the laptop's webcam is the *system's* input camera (what LensaRak actually uses to detect items) — don't mix these two up. Position the laptop so its screen and the item table are both in R4's frame at once, or have R4 pan smoothly between them.

Pre-Roll Checklist

- Check best.pt is the RIGHT model first — the one in the folder now may be an older run, not the latest merged-dataset training. Check runs\detect\train-4\results.csv for the newest completed epoch, then promote it: copy "runs\detect\train-4\weights\best.pt" "best.pt" (safe any time — checkpoints save every epoch, training doesn't need to finish first)
- Run a silent test pass (no recording) — confirm detection works, no errors
- Items arranged: a mix of easy-to-detect items + ONE item deliberately set up to fail
- lensarak_review.py queue is empty before starting
- Laptop charged / plugged in
- Good lighting, tidy background
- All 4 people know their cue — do ONE dry run without recording first

Full Script

[0:00–0:20] Team & Project Intro

R1 (to camera):

"Hi, we're Team LensaRak, from UniKL BMI. This is our submission for the AI Innovators Challenge — a camera-based inventory system for Petrosains' storeroom."

(R2, R3, R4 each say their name in ~2 seconds — quick, not scripted individually, just first name + a nod)

[0:20–0:50] The Problem

R1 (to camera, gesturing to the table):

"Petrosains' storeroom has over a hundred catalogued items used across daily programmes. When a session ends, staff do a bulk return — dumping mixed items back all at once, while the next activation is already loading in. That moment is logged manually today, under time pressure — exactly where miscounts happen. LensaRak automates that moment."

[0:50–1:10] Launching the System

R3 (at laptop): double-clicks LensaRak.bat, selects option 3 — Bulk return, webcam, live

R1 (narrating):

"Our system is camera-agnostic — it runs on a laptop webcam or a phone acting as an IP camera, whichever a store already has. Here we're using the webcam. Let's start a bulk return."

[1:10–2:40] Live Bulk-Return Scan (core demo)

R2: places a mix of items on the table in front of the camera (2–4 items, including one bagged/packaged item if possible)

R3: presses SPACE to scan

R1 (narrating over the detection):

"The system uses YOLOv8 to detect each item and read its confidence score. Anything above 0.75 gets applied to the ledger automatically. Watch the confidence numbers on screen."

R3: confirms quantities as prompted (ENTER to accept, or U to type a real count for a bagged item)

R1 (if an OCR-read label appears):

"This bag has a printed quantity — the system reads it automatically with OCR, so no one has to count by hand."

[2:40–3:40] Deliberate Failure Case

R2: places or angles ONE item so the camera struggles to identify it confidently (partial occlusion, odd angle, or a class the model hasn't seen much)

R3: presses SPACE again

R1 (narrating, pointing at the low confidence number):

"Here the confidence is low. Instead of guessing, the system does NOT update stock — it flags this item for human review. This is deliberate: LensaRak is built to ask a human when it isn't sure, rather than silently getting it wrong."

R3: presses Q to close the live window

[3:40–4:30] Recovery — Human Review

R3 (at laptop): runs the review tool (lensarak_review.py or option 5 in LensaRak.bat)

R1 (narrating):

“This is our review console. It shows exactly what the camera saw and why it was flagged. A reviewer can confirm it’s correct, fix the quantity, correct the item entirely, or reject it — and the system keeps a full record of both the camera’s original read and the human’s correction, so nothing is ever silently overwritten.”

R3: demonstrates one option live (e.g. presses C to correct the item, or ENTER to approve)

[4:30–5:00] Closing

R1 (to camera):

“So LensaRak automates what it can confidently handle, and hands off to a human exactly when it can’t. That’s the core design of our system. Thank you.”

(all 4 face camera briefly, end recording)

Notes

- Keep pacing brisk — rehearse once before the real take so nobody’s reading word-for-word on camera; use the lines above as a guide, not something to recite stiffly.
- If a line runs long and you’re close to 5:00, cut from the closing statement first, not from the failure-case/recovery section — that part is a hard requirement.
- R1 can hold a printed or phone copy of this script off-camera as a cue sheet.